

# PARALLEL IMPLEMENTATION OF MEDIAN FILTERING

## PROJECT REPORT

Name: Prakhar Vipin Jain

Course: MATH 424

Department: CPRE

Net-ID: [pvjain05@iastate.edu](mailto:pvjain05@iastate.edu)

## INTRODUCTION

---

Image filters can be classified as linear or nonlinear. Linear filters are also known as convolution filters as they can be represented using a matrix multiplication. Thresholding and image equalization are examples of nonlinear operations, as is the median filter. Median filtering is a nonlinear operation often used in image processing to reduce "salt and pepper" noise. A median filter is more effective than convolution when the goal is to simultaneously reduce noise and preserve the edges in an image.

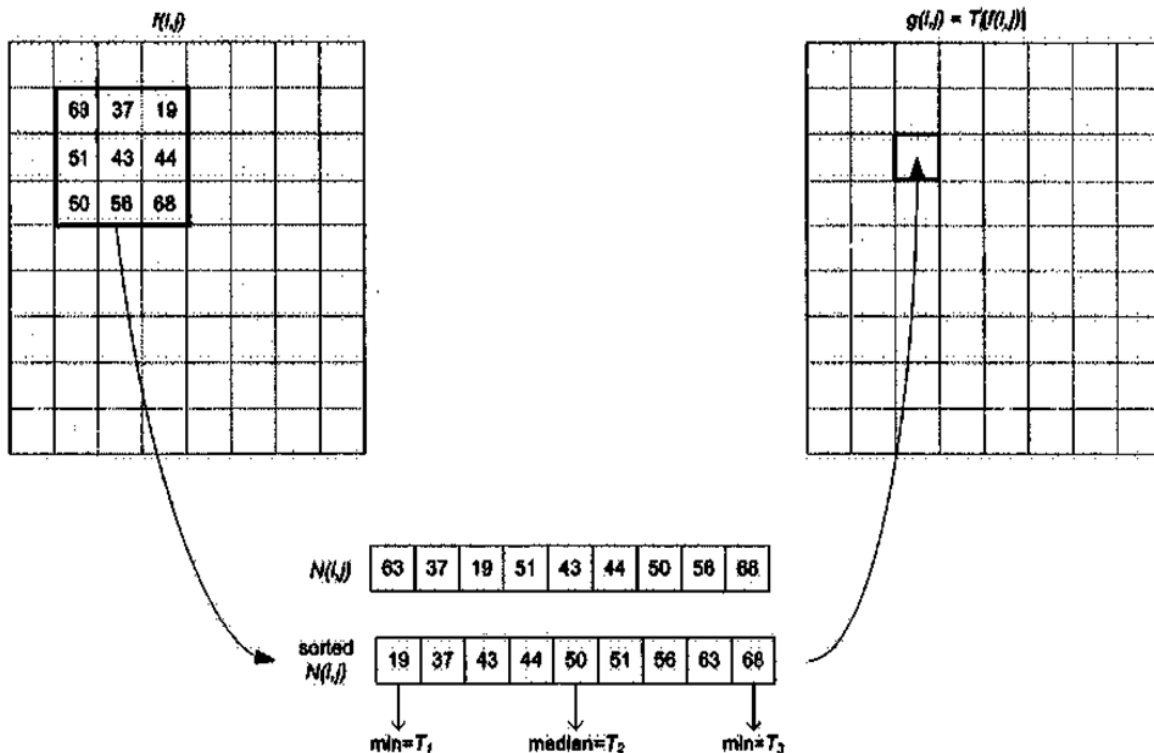
In median filtering, the noisy pixel considered is replaced by the median value of the neighboring pixels. The number of neighboring pixels depend on the size of the mask. Typically, a  $(3 \times 3)$  mask is the smallest mask that can be used. Other than this, mask sizes of  $(5 \times 5)$ ,  $(7 \times 7)$ ,  $(9 \times 9)$  and so on can also be used. Thus, the median filter works by sliding through the image pixel by pixel and replacing them by the median value obtained among the neighboring pixels. The median is calculated by first sorting the values of all the pixels in the increasing order of magnitude and then picking up the middle value from this sorted array of pixel values. This is one of the reasons why the size of the mask is always an odd multiple.

An image consists of millions of pixels and better the quality of the image, more the number of pixels. Processing each and every pixel is a time-consuming task and if done sequentially, i.e. for every mask, sorting the pixel values and then replacing the considered pixel with the median value ultimately results in slow processing of the image. This time is inversely proportional to the size of the mask. This means that as the size of the mask decreases, the processing time increases.

Comparatively, parallelizing the processing task can help reduce the processing time of the image considerably. The serial program for the above implementation of the median filter has been constructed and is provided later in this report. Looking ahead, I will parallelize the deduced serial code in OpenMP. This will include dividing the iterations pertaining to the sorting the pixel values and finding the median among them for every sliding window.

## WORKED OUT EXAMPLE

Given below is an example of how Median Filtering is done on an image. Consider the matrix given below as the part of the image. Here a (3 x 3) window is used to filter the given image. For altering the boundary elements of the image, their duplicates are considered at the respective edge. The elements or the pixel values within the window size are sorted and the considered pixel is replaced with the median value.



## THE SELECTED OPENMP CLAUSES

---

```
# pragma omp parallel for collapse(2) ordered num_threads(thread_count)
default(none) shared(image, filteredimage, rows, cols, n, time)
firstprivate(pixel_values) private(my_rank, t3, t4, p, rr, cc)
schedule(static,1)
```

The above schedule clause which is 'static' with chunk size 1 is selected because it has the best performance among all the other scheduling techniques which included dynamic, guided, auto and runtime. Along with running the program with different schedules, I also tried different schedules with different chunk sizes, viz. 1, 4, 8, 16 and default chunk sizes.

For static schedule, the default value of the chunk size is the number of iterations divided by the number of threads i.e.  $\frac{\text{Number of iterations}}{\text{Number of threads}}$ . For dynamic and guided schedules, the default value for the chunk size is 1. As for auto and runtime schedules, there is no chunk size. The compiler allocates the threads at runtime.

As the section of the code that I have parallelized contains 'for' loops, I have used 'parallel for' construct for parallelization. I have used an additional collapse clause which I will mention later. Another construct I have used is the 'ordered' construct which is used to sequentialize and orders the execution of ordered regions. Unordered populating of the rows and columns with the filtered pixel values would have distorted the image. Thus, I have used the 'ordered' construct while populating the rows and columns of the output image.

The scope of the variables was also defined as for parallelization, some variables needed to be shared, private and firstprivate. The variables that were shared are {image, filteredimage, rows, cols, n, time}. The variable 'image' is a 2D vector which contains the input image. As each thread is only reading the image, there is no problem of race conditions between the threads. The 'filteredimage' variable is a pointer to a 2D vector and points to the output image. It is shared among the thread because it is the output image and every thread has to populate the rows and columns of it by the pixels they have computed individually. The variables, 'rows' and 'columns' contains the number of rows and columns of the input image passed and 'n' is the size of the window/mask that performs the median filtering over (n x n) area in the image. These variables are also shared among the

threads because the threads only have to read them and so there is no problem of race conditions. The last shared variable is 'time' which is an array used to store the time required by each thread to compute their own chunk. As each thread stores their time in it based on their rank, this variable is also shared.

The variables that were private to each thread are {pixel\_values, my\_rank, t3, t4, p, rr, cc}. The variable 'pixel\_values' is first private because it needed to be initialized with the value that it encounters in the previous construct. If it would be private, every time the thread returned after an iteration, the value of pixel\_values would be zero which we do not want. Other variables like 'my\_rank', 't3' and 't4' relate to an individual thread which is reason for them to be private. Variables 'p', 'rr' and 'cc' are used by each thread to compute their own chunk and so are also private.

## OPENMP NEW CONCEPT USED

---

The collapse clause may be used to specify how many loops are associated with the loop construct. The parameter of the collapse clause must be a constant positive integer expression. If a collapse clause is specified with a parameter value greater than 1, then the iterations of the associated loops to which the clause applies are collapsed into one larger iteration space that is then divided according to the schedule clause. The sequential execution of the iterations in these associated loops determines the order of the iterations in the collapsed iteration space. If no collapse clause is present or its parameter is, the only loop that is associated with the loop construct for the purposes of determining how the iteration space is divided according to the schedule clause is the one that immediately follows the loop directive.

I have mentioned before that I have used another clause named 'collapse' which is used to increase the total number of iterations that will be partitioned across the available number of OMP threads by reducing the granularity of work to be done by each thread. It also increases the scalability of the program. In this program, I have collapsed two for loops (i and j) thereby increasing the iteration space for rows and columns of the image.

## RESULTS

---

The input and output images for both serial and parallel execution are of .pgm format which is a Portable Grayscale Map. A PGM image represents a grayscale graphic image. There are many pseudo-PGM formats in use where everything is as specified herein except for the meaning of individual pixel values. For most purposes, a PGM image can just be thought of an array of arbitrary integers, and all the programs in the world that think they're processing a grayscale image can easily be tricked into processing something else.

\*The results displayed below are for Input\_Image3.pgm

**For serial execution –**

Input Image:

Output Image:



This above filtering is done by using window size = 9. One can notice the stars to be blur in the output image which indicates reduction of noise from the image.

Execution time = 263.193462 secs

**For parallel execution –**

Input Image:



Output Image:



This is the output image with static scheduling with chunk size 1 which proved to be the optimum schedule for parallelization.

Execution time = 92.773567 secs

**Verification with MATLAB-**





# DATA ANALYSIS

**\*NOTE- All timings are in seconds**

**Data Set 1: Input\_Image1.pgm (18,994 KB)**

| Threads<br>→             | Chunk<br>size<br>↓ | 1          | 2          | 4          | 8          | 16         |
|--------------------------|--------------------|------------|------------|------------|------------|------------|
| Schedules:<br><br>Static | 1                  | 159.845610 | 117.618669 | 87.075414  | 73.088391  | 59.778781  |
|                          | 4                  | 159.741286 | 132.329936 | 127.275453 | 147.192496 | 156.108165 |
|                          | 8                  | 159.994710 | 147.171165 | 144.139437 | 163.814115 | 171.852579 |
|                          | 16                 | 160.001472 | 153.559104 | 153.040254 | 171.760347 | 205.658922 |
| Dynamic                  | 1                  | 159.870233 | 115.997256 | 87.090787  | 72.831194  | 60.253612  |
|                          | 4                  | 160.067111 | 132.131750 | 127.503150 | 149.048226 | 152.275864 |
|                          | 8                  | 160.124409 | 146.360010 | 145.586016 | 162.689667 | 170.716539 |
|                          | 16                 | 160.017528 | 153.391619 | 153.119732 | 172.553518 | 203.815412 |
| Guided                   | 1                  | 160.182147 | 160.649851 | 160.002478 | 160.137532 | 162.481793 |
|                          | 4                  | 160.065258 | 161.689344 | 160.550149 | 160.199708 | 160.618009 |
|                          | 8                  | 160.112786 | 161.259638 | 160.944127 | 161.257453 | 160.583052 |
|                          | 16                 | 159.954120 | 160.285715 | 160.129530 | 160.128016 | 160.743391 |
| Auto                     | No chunk<br>size   | 159.585215 | 160.070749 | 160.068013 | 161.014960 | 160.257501 |
| Runtime                  |                    | 159.844471 | 103.295304 | 75.054813  | 74.117641  | 60.477173  |

- The serial time = 160.820007 secs
- Best case Parallel execution time = 59.778781 secs
- The workload distribution for static schedule was even and it is clear from the execution times of each thread out of 16 threads given below:

*Time taken by thread 0 to compute its chunk = 59.175427 secs*

*Time taken by thread 1 to compute its chunk = 59.308516 secs*

*Time taken by thread 2 to compute its chunk = 59.085974 secs*

*Time taken by thread 3 to compute its chunk = 59.169525 secs*

*Time taken by thread 4 to compute its chunk = 59.005827 secs*

*Time taken by thread 5 to compute its chunk = 59.196315 secs*

*Time taken by thread 6 to compute its chunk = 59.231253 secs*

*Time taken by thread 7 to compute its chunk = 59.206661 secs*

*Time taken by thread 8 to compute its chunk = 59.226837 secs*

*Time taken by thread 9 to compute its chunk = 59.045139 secs*

*Time taken by thread 10 to compute its chunk = 58.899492 secs*

*Time taken by thread 11 to compute its chunk = 59.053454 secs*

*Time taken by thread 12 to compute its chunk = 58.921677 secs*

*Time taken by thread 13 to compute its chunk = 58.893360 secs*

*Time taken by thread 14 to compute its chunk = 58.924835 secs*

*Time taken by thread 15 to compute its chunk = 58.895725 secs*

- The auto and guided schedule show similar timings that is fairly close to the serial execution time. This may be because of the fact that guided schedule allots a chunk of data to the threads only when the thread is done computing its previous chunk and asks for a new one. But as the parallel code contains a critical as well as an ordered clause, there is lot of overhead as threads have to wait for a long time along with their chunks to be executed.

The workload distribution was uneven for the 'auto' schedule:

*Time taken by thread 0 to compute its chunk = 10.605937 secs*

*Time taken by thread 1 to compute its chunk = 21.073863 secs*

*Time taken by thread 2 to compute its chunk = 31.768825 secs*

*Time taken by thread 3 to compute its chunk = 42.804017 secs*

*Time taken by thread 4 to compute its chunk = 54.107703 secs*

*Time taken by thread 5 to compute its chunk = 65.698848 secs*

*Time taken by thread 6 to compute its chunk = 77.231345 secs*

*Time taken by thread 7 to compute its chunk = 88.536510 secs*

*Time taken by thread 8 to compute its chunk = 100.280961 secs*

*Time taken by thread 9 to compute its chunk = 111.037032 secs*

*Time taken by thread 10 to compute its chunk = 120.036634 secs*

*Time taken by thread 11 to compute its chunk = 128.396288 secs*

*Time taken by thread 12 to compute its chunk = 135.420295 secs*

*Time taken by thread 13 to compute its chunk = 143.372302 secs*

*Time taken by thread 14 to compute its chunk = 152.289238 secs*

*Time taken by thread 15 to compute its chunk = 160.213678 secs*

- **This uneven distribution was true for 'auto' schedule for all the other data sets as well.**
- I tried running the best schedule i.e. static with chunk size 1 with "numactl --membind=0,1 ./median\_parallel.sh". The execution time was **59.006019** secs which is further less than 59.778781 secs.
- As the execution time for parallel execution started increasing from the 16<sup>th</sup> thread for chunk size 1, I did not calculate for higher number of threads.

**Data Set 2: Input\_Image2.pgm (23,438 KB)**

| Threads<br>→             | Chunk<br>size<br>↓ | 1          | 2          | 4          | 8          | 16         |
|--------------------------|--------------------|------------|------------|------------|------------|------------|
| Schedules:<br><br>Static | 1                  | 202.868598 | 155.673198 | 107.702861 | 97.195866  | 74.463772  |
|                          | 4                  | 202.218307 | 167.885854 | 163.828961 | 182.969107 | 195.442512 |
|                          | 8                  | 202.468531 | 185.484520 | 176.545186 | 198.018291 | 209.739152 |
|                          | 16                 | 202.451368 | 189.788451 | 189.017553 | 210.545218 | 238.048461 |
| Dynamic                  | 1                  | 202.054862 | 149.843054 | 108.030969 | 93.826356  | 75.680126  |
|                          | 4                  | 202.424351 | 136.485183 | 119.554039 | 143.704134 | 161.065130 |
|                          | 8                  | 201.878542 | 145.184465 | 144.845159 | 198.516521 | 210.884027 |
|                          | 16                 | 201.715380 | 153.498456 | 153.197602 | 209.148539 | 228.541498 |
| Guided                   | 1                  | 201.985451 | 202.763191 | 202.193746 | 202.784561 | 202.681002 |
|                          | 4                  | 202.775829 | 202.661943 | 202.200147 | 201.624889 | 202.018359 |
|                          | 8                  | 202.488960 | 201.980238 | 202.605571 | 201.129203 | 201.997420 |
|                          | 16                 | 203.008487 | 202.484611 | 201.958412 | 201.731948 | 201.488351 |
| Auto                     | No chunk<br>size   | 202.487206 | 201.991024 | 202.624801 | 202.485100 | 202.548015 |
| Runtime                  |                    | 202.668514 | 123.947124 | 111.808369 | 93.687449  | 76.087816  |

- The serial time = 204.369995 secs
- Best case Parallel execution time = 74.463772 secs
- The workload distribution for static schedule was even and it is clear from the execution times of each thread out of 16 threads given below:

*Time taken by thread 0 to compute its chunk = 74.989118 secs*

*Time taken by thread 1 to compute its chunk = 74.854582 secs*

*Time taken by thread 2 to compute its chunk = 74.757406 secs*

*Time taken by thread 3 to compute its chunk = 74.926200 secs*

*Time taken by thread 4 to compute its chunk = 74.896859 secs*

*Time taken by thread 5 to compute its chunk = 74.953371 secs*

*Time taken by thread 6 to compute its chunk = 75.080892 secs*

*Time taken by thread 7 to compute its chunk = 75.020143 secs*

*Time taken by thread 8 to compute its chunk = 74.917979 secs*

*Time taken by thread 9 to compute its chunk = 74.909147 secs*

*Time taken by thread 10 to compute its chunk = 74.639726 secs*

*Time taken by thread 11 to compute its chunk = 74.650440 secs*

*Time taken by thread 12 to compute its chunk = 74.752266 secs*

*Time taken by thread 13 to compute its chunk = 74.680941 secs*

*Time taken by thread 14 to compute its chunk = 74.678205 secs*

*Time taken by thread 15 to compute its chunk = 74.673349 secs*

- I tried running the best schedule i.e. static with chunk size 1 with “numactl --membind=0,1 ./median\_parallel.sh”. The execution time was **74.219392** secs which is further less than 74.463772 secs.

**Data Set 3: Input\_Image3.pgm (27,547 KB)**

| Threads<br>→ | Chunk<br>size<br>↓ | 1          | 2          | 4          | 8          | 16         |
|--------------|--------------------|------------|------------|------------|------------|------------|
| Schedules:   | 1                  | 263.193462 | 173.573315 | 135.101077 | 117.154057 | 92.773567  |
|              | 2                  | 262.936057 | 173.650104 | 148.154255 | 168.441251 | 178.032077 |
|              | 4                  | 263.778588 | 215.011170 | 208.020343 | 236.099375 | 249.875132 |
|              | 8                  | 263.395564 | 239.524179 | 236.395383 | 268.244220 | 280.195565 |
|              | 16                 | 263.786110 | 251.829353 | 251.352660 | 282.789761 | 338.788450 |
|              | Default            | 263.392294 | 263.720043 | 263.095824 | 263.491843 | 263.631292 |
| Dynamic      | 1                  | 264.383264 | 155.191040 | 137.229584 | 114.507565 | 93.079510  |
|              | 2                  | 264.421474 | 163.394685 | 151.830979 | 164.835083 | 171.337266 |
|              | 4                  | 263.472723 | 211.308588 | 209.950337 | 236.196474 | 248.357794 |
|              | 8                  | 263.684505 | 240.456126 | 236.932889 | 267.279579 | 278.945360 |
|              | 16                 | 263.006030 | 251.670132 | 251.147524 | 281.705669 | 336.945297 |
| Guided       | 1                  | 263.351812 | 263.725511 | 262.886220 | 263.437402 | 263.538888 |
|              | 2                  | 262.935334 | 263.807951 | 263.845951 | 262.584845 | 263.446711 |
|              | 4                  | 263.865233 | 263.771412 | 262.555214 | 263.110249 | 262.473829 |
|              | 8                  | 262.478510 | 263.018544 | 263.885268 | 263.257841 | 263.985177 |
|              | 16                 | 262.110218 | 263.784147 | 263.101128 | 263.124781 | 262.577452 |
| Auto         | No chunk           | 261.665263 | 263.798283 | 262.984063 | 263.504336 | 262.478512 |
| Runtime      | size               | 264.556103 | 209.147851 | 176.001476 | 120.785421 | 94.682036  |

- The serial time = 263.193462 secs
- Best case Parallel execution time = 92.773567 secs
- The workload distribution for static schedule was even and it is clear from the execution times of each thread out of 16 threads given below:

*Time taken by thread 0 to compute its chunk = 92.373467 secs*

*Time taken by thread 1 to compute its chunk = 92.574180 secs*

*Time taken by thread 2 to compute its chunk = 92.641522 secs*

*Time taken by thread 3 to compute its chunk = 92.522662 secs*

*Time taken by thread 4 to compute its chunk = 92.512363 secs*

*Time taken by thread 5 to compute its chunk = 92.570308 secs*

*Time taken by thread 6 to compute its chunk = 92.562178 secs*

*Time taken by thread 7 to compute its chunk = 92.545392 secs*

Time taken by thread 8 to compute its chunk = 92.533315 secs  
Time taken by thread 9 to compute its chunk = 92.658345 secs  
Time taken by thread 10 to compute its chunk = 92.488986 secs  
Time taken by thread 11 to compute its chunk = 92.150117 secs  
Time taken by thread 12 to compute its chunk = 92.200594 secs  
Time taken by thread 13 to compute its chunk = 92.132365 secs  
Time taken by thread 14 to compute its chunk = 92.206376 secs  
Time taken by thread 15 to compute its chunk = 92.359847 secs

- I tried running the best schedule i.e. static with chunk size 1 with “numactl --membind=1 ./median\_parallel.sh”. The execution time was **91.652204** secs which is further less than 92.773567 secs.

#### Data Set 4: Input\_Image4.pgm (40,573 KB)

| Threads<br>→             | Chunk<br>size<br>↓ | 1          | 2          | 4          | 8          | 16         |
|--------------------------|--------------------|------------|------------|------------|------------|------------|
| Schedules:<br><br>Static | 1                  | 320.924301 | 203.797932 | 168.335727 | 154.422265 | 122.814190 |
|                          | 4                  | 319.014128 | 265.834311 | 258.315386 | 297.454754 | 306.212876 |
|                          | 8                  | 319.814752 | 293.499979 | 290.561753 | 325.366131 | 341.880183 |
|                          | 16                 | 319.854721 | 307.311527 | 315.986562 | 347.657032 | 395.676989 |
| Dynamic                  | 1                  | 319.065278 | 205.469308 | 180.126757 | 150.502692 | 125.166953 |
|                          | 4                  | 319.743651 | 259.147529 | 238.875028 | 294.217590 | 309.305614 |
|                          | 8                  | 319.558102 | 267.079641 | 261.066913 | 302.548642 | 323.647100 |
|                          | 16                 | 318.001423 | 290.914560 | 289.156977 | 315.625891 | 336.078178 |
| Guided                   | 1                  | 320.904215 | 321.011538 | 321.001472 | 320.802146 | 320.512184 |
|                          | 4                  | 319.421397 | 319.514489 | 320.744121 | 321.014792 | 321.110457 |
|                          | 8                  | 319.234810 | 320.021549 | 319.874842 | 320.917356 | 321.014736 |
|                          | 16                 | 318.079124 | 321.147260 | 320.904215 | 319.704681 | 320.731841 |
| Auto                     | No                 | 318.661973 | 319.029364 | 319.196331 | 318.510953 | 318.759015 |
| Runtime                  | chunk<br>size      | 320.187503 |            |            |            |            |

- The serial time = 322.440002 secs
- Best case Parallel execution time = 122.814190 secs
- The workload distribution for static schedule was even and it is clear from the execution times of each thread out of 16 threads given below:

Time taken by thread 0 to compute its chunk = 122.228728 secs  
Time taken by thread 1 to compute its chunk = 121.981067 secs  
Time taken by thread 2 to compute its chunk = 122.079668 secs  
Time taken by thread 3 to compute its chunk = 122.142349 secs

Time taken by thread 4 to compute its chunk = 122.233528 secs  
Time taken by thread 5 to compute its chunk = 122.263985 secs  
Time taken by thread 6 to compute its chunk = 122.359473 secs  
Time taken by thread 7 to compute its chunk = 122.300504 secs  
Time taken by thread 8 to compute its chunk = 122.063045 secs  
Time taken by thread 9 to compute its chunk = 122.175282 secs  
Time taken by thread 10 to compute its chunk = 121.614302 secs  
Time taken by thread 11 to compute its chunk = 121.623304 secs  
Time taken by thread 12 to compute its chunk = 121.514588 secs  
Time taken by thread 13 to compute its chunk = 121.604645 secs  
Time taken by thread 14 to compute its chunk = 121.675647 secs  
Time taken by thread 15 to compute its chunk = 121.569140 secs

- I tried running the best schedule i.e. static with chunk size 1 with “numactl --membind=1 ./median\_parallel.sh”. The execution time was **121.315165**secs which is further less than 122.814190 secs.

#### **Data Set 5: Input\_Image5.pgm (41,189 KB)**

| Threads<br>→             | Chunk<br>size<br>↓ | 1          | 2          | 4          | 8          | 16         |
|--------------------------|--------------------|------------|------------|------------|------------|------------|
| Schedules:<br><br>Static | 1                  | 317.307204 | 221.980497 | 171.382066 | 149.577882 | 122.572143 |
|                          | 4                  | 316.736039 | 262.641485 | 252.715101 | 292.006935 | 304.854894 |
|                          | 8                  | 316.041560 | 289.860343 | 285.962997 | 321.478609 | 338.540281 |
|                          | 16                 | 317.114723 | 306.500580 | 303.485367 | 340.555992 | 399.463339 |
|                          | Default            | 318.349357 | 317.181895 | 316.568222 | 316.785780 | 316.977620 |
| Dynamic                  | 1                  | 318.684126 | 235.159015 | 169.650887 | 149.213248 | 123.151026 |
|                          | 4                  | 318.741452 | 261.880632 | 254.220357 | 292.812627 | 300.445066 |
|                          | 8                  | 321.844758 | 289.923953 | 286.936467 | 323.372678 | 348.505445 |
|                          | 16                 | 317.416015 | 303.859942 | 303.520297 | 352.360554 | 400.940202 |
| Guided                   | 1                  | 317.236978 | 316.847520 | 317.235275 | 317.500143 | 317.040865 |
|                          | 4                  | 316.147520 | 317.199687 | 316.774920 | 316.845112 | 317.107772 |
|                          | 8                  | 316.114759 | 318.222753 | 317.964102 | 316.674102 | 316.112473 |
|                          | 16                 | 317.521470 | 316.740023 | 316.680122 | 317.657012 | 318.831931 |
| Auto                     | No chunk           | 317.171288 | 316.815879 | 316.605792 | 317.761608 | 316.632227 |
| Runtime                  | size               | 316.632227 | 198.991419 | 171.962305 | 147.412915 | 123.885870 |

- The serial time = 318.700012 secs
- Best case Parallel execution time = 122.572143 secs

- The workload distribution for static schedule was even and it is clear from the execution times of each thread out of 16 threads given below:

*Time taken by thread 0 to compute its chunk = 121.733547 secs*

*Time taken by thread 1 to compute its chunk = 122.140164 secs*

*Time taken by thread 2 to compute its chunk = 122.249230 secs*

*Time taken by thread 3 to compute its chunk = 121.971068 secs*

*Time taken by thread 4 to compute its chunk = 121.959221 secs*

*Time taken by thread 5 to compute its chunk = 121.867586 secs*

*Time taken by thread 6 to compute its chunk = 122.090898 secs*

*Time taken by thread 7 to compute its chunk = 122.162658 secs*

*Time taken by thread 8 to compute its chunk = 122.206198 secs*

*Time taken by thread 9 to compute its chunk = 122.036070 secs*

*Time taken by thread 10 to compute its chunk = 121.479819 secs*

*Time taken by thread 11 to compute its chunk = 121.435774 secs*

*Time taken by thread 12 to compute its chunk = 121.574673 secs*

*Time taken by thread 13 to compute its chunk = 121.449644 secs*

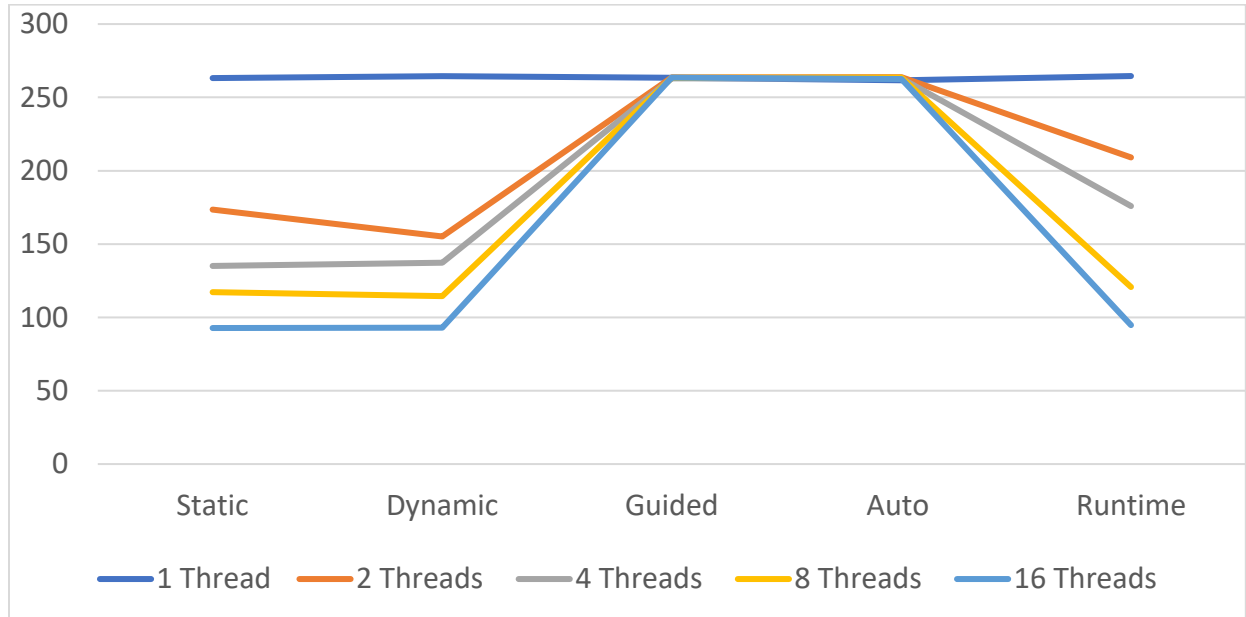
*Time taken by thread 14 to compute its chunk = 121.457294 secs*

*Time taken by thread 15 to compute its chunk = 121.275871 secs*

- I tried running the best schedule i.e. static with chunk size 1 with “numactl --membind=0,1 ./median\_parallel.sh”. The execution time was **122.101921** secs which is further less than 122.572143 secs.

# STATISTICAL ANALYSIS

Execution time vs Schedules for different number of threads for Input\_Image3.pgm-



The above trend is the same for all the 5 output images

Serial Execution Time vs Best-case Parallel Execution Time for all images-





# PERFORMANCE ANALYSIS

---

## Data Set 1: Input\_Image1.pgm (18,994 KB)

$$\begin{aligned}\text{Speedup} &= \frac{T_{\text{serial}}}{T_{\text{parallel}}} \\ &= \frac{160.820007}{59.778781} \\ &\sim 2.69\end{aligned}$$

$$\begin{aligned}\text{Efficiency} &= \frac{\text{Speedup}}{\text{No.of threads}} \\ &= \frac{2.69}{16} \\ &= 16.8\%\end{aligned}$$

## Data Set 2: Input\_Image2.pgm (23,438 KB)

$$\begin{aligned}\text{Speedup} &= \frac{T_{\text{serial}}}{T_{\text{parallel}}} \\ &= \frac{204.369995}{74.463772} \\ &\sim 2.74\end{aligned}$$

$$\begin{aligned}\text{Efficiency} &= \frac{\text{Speedup}}{\text{No.of threads}} \\ &= \frac{2.74}{16} \\ &= 17.12\%\end{aligned}$$

## Data Set 3: Input\_Image3.pgm (27,547 KB)

$$\begin{aligned}\text{Speedup} &= \frac{T_{\text{serial}}}{T_{\text{parallel}}} \\ &= \frac{263.193462}{92.773567} \\ &\sim 2.83\end{aligned}$$

$$\begin{aligned}\text{Efficiency} &= \frac{\text{Speedup}}{\text{No.of threads}} \\ &= \frac{2.83}{16} \\ &= 17.5\%\end{aligned}$$

#### **Data Set 4: Input\_Image4.pgm (40,573 KB)**

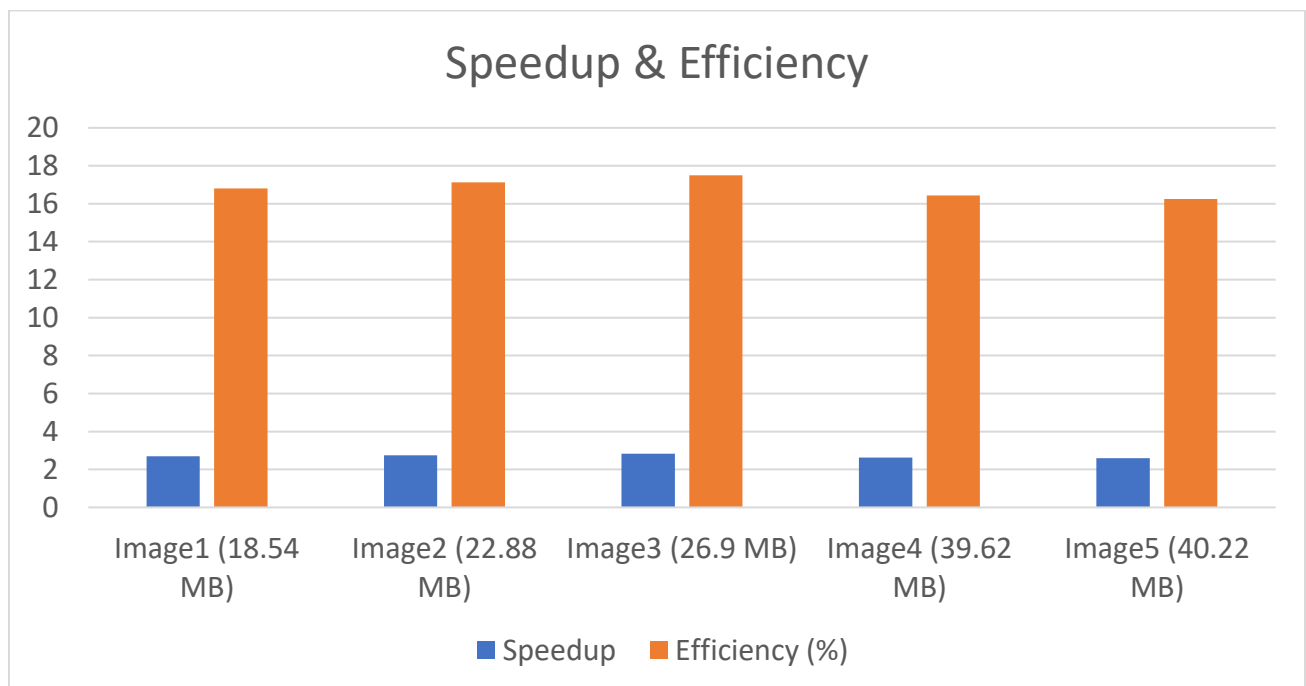
$$\begin{aligned}\text{Speedup} &= \frac{T_{\text{serial}}}{T_{\text{parallel}}} \\ &= \frac{322.440002}{122.814190} \\ &\sim 2.63\end{aligned}$$

$$\begin{aligned}\text{Efficiency} &= \frac{\text{Speedup}}{\text{No. of threads}} \\ &= \frac{2.63}{16} \\ &= 16.43\%\end{aligned}$$

#### **Data Set 5: Input\_Image5.pgm (41,189 KB)**

$$\begin{aligned}\text{Speedup} &= \frac{T_{\text{serial}}}{T_{\text{parallel}}} \\ &= \frac{318.700012}{122.572143} \\ &\sim 2.6\end{aligned}$$

$$\begin{aligned}\text{Efficiency} &= \frac{\text{Speedup}}{\text{No. of threads}} \\ &= \frac{2.6}{16} \\ &= 16.25\%\end{aligned}$$



**Scalability:**

- We can see that the efficiency is almost constant for every data set and there is no significant variation as the image size increases.
- That is, the efficiency does not change as the data set/size increases.
- Thus, the parallelized program can be said to be weakly scalable.

# APPENDIX

---

## A. Parallel Code:

There are five files included: The main program c file, a function c file, a header file and two shell files for running the serial and parallel programs.

### 1. median\_parallel.c

```
/*
*****
* Program: median_parallel.c
* Purpose: This program will apply a median filter to an image with a user
* specified window size.
* Run: ./median_parallel.sh
* Name: Prakhar Jain, MATH 424
*****
**/
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <time.h>
#include <omp.h>
#include "imageio.h"

void insertion_sort(unsigned char *item, int count);

int main(int argc, char *argv[])
{
    char *inputfilename1=NULL, *inputfilename2=NULL, *inputfilename3=NULL,
    *inputfilename4=NULL, *inputfilename5=NULL;
    char *outputfilename = NULL;
    int n=0, rows, cols, i, c, flag=0;
    unsigned char **image1=NULL, **image2=NULL, **image3=NULL, **image4=NULL,
    **image5=NULL, **filteredimage1 = NULL, **filteredimage2 = NULL,
    **filteredimage3 = NULL, **filteredimage4 = NULL, **filteredimage5 = NULL;
    int thread_count,t;
    double t1,t2;
    void median_filter(unsigned char **image, int rows, int cols, int n,
    unsigned char ***filteredimage, int thread_count);

    /*
    *****
    * Get the command line parameters.
    *****
    */
    for(i=1;i<argc;i++)
    {
```

```

        if(strcmp(argv[i], "-n") == 0)
        {
            n = atoi(argv[i+1]);
            i++;
        }
        else
        {
            if(inputfilename1 == NULL) inputfilename1 = argv[i];
            else if(inputfilename2 == NULL) inputfilename2 = argv[i];
            else if(inputfilename3 == NULL) inputfilename3 = argv[i];
            else if(inputfilename4 == NULL) inputfilename4 = argv[i];
            else if(inputfilename5 == NULL) inputfilename5 = argv[i];
            else if(outputfilename == NULL) outputfilename = argv[i];
            else thread_count = strtol(argv[i], NULL, 10);
        }
    }

    if((n <= 0) || (inputfilename1==NULL) || (inputfilename2==NULL) ||
(inputfilename3==NULL) || (inputfilename4==NULL) || (inputfilename5==NULL)
||(outputfilename==NULL) || (thread_count <= 0))
    {

printf("\n*****\n");
        printf("This program will apply a median filter to an image. You need
to specify the\n");
        printf("size of the window to use in median filtering the image (i.e. -
n 5), \n");
        printf("the 5-input images to process, the name of an output file\n");
        printf("in which the median filtered image will be written and the
number of threads.\n");
        printf("Thus, you could run the program as follows:\n");
        printf("\n");
        printf("Ex: median -n 9 Input_Image1.pgm Input_Image2.pgm
Input_Image3.pgm Input_Image4.pgm Input_Image5.pgm medianfiltered.pgm 8\n");
        printf("\n");

printf("*****\n");

        fprintf(stderr, "\n<USAGE> median -n # inputPGMfile outputPGMfile
#of_threads\n\n");
        exit(1);
    }

/*****
    * printf("Inputfilename = %s\n", inputfilename);
    * printf("Outputfilename = %s\n", outputfilename);

```

```

    * printf("n = %d\n", n);
    *****/

/*****
    * Read in the PGM image from the file.

    *****/

    printf("For Input_Image1.pgm - 1\n");
    printf("For Input_Image2.pgm - 2\n");
    printf("For Input_Image3.pgm - 3\n");
    printf("For Input_Image4.pgm - 4\n");
    printf("For Input_Image5.pgm - 5\n");
    printf("Choose which image you want to smooth using median filter : ");
    scanf("%d", &c);
    printf("\n");

    switch(c)
    {
        case 1:
        {
            if(read_pgm_image(inputfilename1, &image1, &rows, &cols) == 0)
            exit(1);
            flag=1;
        }
        break;
        case 2:
        {
            if(read_pgm_image(inputfilename2, &image2, &rows, &cols) == 0)
            exit(1);
            flag=2;
        }
        break;
        case 3:
        {
            if(read_pgm_image(inputfilename3, &image3, &rows, &cols) == 0)
            exit(1);
            flag=3;
        }
        break;
        case 4:
        {
            if(read_pgm_image(inputfilename4, &image4, &rows, &cols) == 0)
            exit(1);
            flag=4;
        }
        break;
        case 5:
        {

```

```

        if(read_pgm_image(inputfilename5, &image5, &rows, &cols) == 0)
exit(1);
        flag=5;
    }
    break;
}

/*****
 * Median filter the image.
*****/

if(flag==1)
{
    t1= omp_get_wtime();
    median_filter(image1, rows, cols, n, &filteredimage1, thread_count);
    t2= omp_get_wtime();
}
if(flag==2)
{
    t1= omp_get_wtime();
    median_filter(image2, rows, cols, n, &filteredimage2, thread_count);
    t2= omp_get_wtime();
}
if(flag==3)
{
    t1= omp_get_wtime();
    median_filter(image3, rows, cols, n, &filteredimage3, thread_count);
    t2= omp_get_wtime();
}
if(flag==4)
{
    t1= omp_get_wtime();
    median_filter(image4, rows, cols, n, &filteredimage4, thread_count);
    t2= omp_get_wtime();
}
if(flag==5)
{
    t1= omp_get_wtime();
    median_filter(image5, rows, cols, n, &filteredimage5, thread_count);
    t2= omp_get_wtime();
}

/*****
 * Print out the size of window used to compute the median
 * and the time taken to do the median filtering.
*****/

```

```

*****/
    printf("Window size = %d (You can increase the window size (only odd
numbers) for more smoothing of the image) \nTime taken by entire parallel
region for median filtering = %lf \n", n, (t2-t1));

    printf("Check the local folder for the MedianFiltered_Image.pgm \n");

/*****
    * Write the filtered image out to a file.
*****/

if(flag==1)
{
    if((write_pgm_image(outputfilename, filteredimage1, rows, cols,
(unsigned char *)NULL, 255)) == 0)
        exit(1);
    free_image(image1, rows);
    free_image(filteredimage1, rows);
}
if(flag==2)
{
    if((write_pgm_image(outputfilename, filteredimage2, rows, cols,
(unsigned char *)NULL, 255)) == 0)
        exit(1);
    free_image(image2, rows);
    free_image(filteredimage2, rows);
}
if(flag==3)
{
    if((write_pgm_image(outputfilename, filteredimage3, rows, cols,
(unsigned char *)NULL, 255)) == 0)
        exit(1);
    free_image(image3, rows);
    free_image(filteredimage3, rows);
}
if(flag==4)
{
    if((write_pgm_image(outputfilename, filteredimage4, rows, cols,
(unsigned char *)NULL, 255)) == 0)
        exit(1);
    free_image(image4, rows);
    free_image(filteredimage4, rows);
}
if(flag==5)
{
    if((write_pgm_image(outputfilename, filteredimage5, rows, cols,
(unsigned char *)NULL, 255)) == 0)

```



```

        exit(1);
        free_image(image5, rows);
        free_image(filteredimage5, rows);
    }

}

/*****
***
* Function: median_filter
* Purpose: This function will median filter an image using an n x n window.
*****/
**/
void median_filter(unsigned char **image, int rows, int cols, int n, unsigned
char ***filteredimage, int thread_count)
{
    unsigned char *pixel_values=NULL;
    int r,t, c, rr, cc, p,my_rank;
    double t3,t4;
    double *time = (double*) malloc((thread_count)*sizeof(double));
    for(t=0;t<thread_count;t++)
        time[t] = 0.0;

/*****
* Allocate an array to store pixel values. There will be up to n x n pixel
* values to sort at each pixel location in the image.
*****/
    if((pixel_values = (unsigned char *) malloc((n*n) * sizeof(unsigned
char))) == NULL){
        fprintf(stderr, "Error allocating an array in median_filter().\n");
        exit(1);
    }

/*****
* Allocate the filtered image.
*****/
    if((*filteredimage) = allocate_image(rows, cols)) == NULL) exit(1);

/*****
* Scan through the image and compute the median of the local pixel values
* at each pixel position.
*****/

```

```

    # pragma omp parallel for collapse(2) ordered num_threads(thread_count)
    default(none) shared(image, filteredimage, rows,cols,n,time)
    firstprivate(pixel_values) private(my_rank,t3,t4,p,rr,cc) schedule(dynamic,1)
    for(r=0;r<rows;r++)
    {
        for(c=0;c<cols;c++)
        {
            my_rank = omp_get_thread_num();
            t3 = omp_get_wtime();
            p=0;
            for(rr=(r-(n/2));rr<(r-(n/2)+n);rr++)
            {
                for(cc=(c-(n/2));cc<(c-(n/2)+n);cc++)
                {
                    if((rr>=0)&&(rr<rows)&&(cc>=0)&&(cc<cols))
                    {
                        pixel_values[p] = image[rr][cc];
                        p++;
                    }
                }
            }
        }
    }

    /*****
    * Sort the array of pixels. Although there can be up
    * to n x n pixels in the array, there are actually only p values.
    *****/
    #pragma omp critical
    {
        insertion_sort(pixel_values, p);
    }

    /*****
    * Assign the median pixel value to the filtered image.
    *****/
    #pragma omp ordered
    {
        (*filteredimage)[r][c] = pixel_values[p/2];
    }
    t4= omp_get_wtime();
    time[my_rank]+=(t4-t3);
}

}
for(t=0;t<thread_count;t++)
    printf("Time taken by thread %d to compute its chunk = %lf\n",t,time[t]);
    free(pixel_values);
}

```

```

void insertion_sort(unsigned char *item, int count)
{
    int c,d,t;
    for (c = 1 ; c <= count; c++)
    {
        d = c;
        while ( d > 0 && item[d-1] > item[d])
        {
            t = item[d];
            item[d] = item[d-1];
            item[d-1] = t;
            d--;
        }
    }
}

```

## 2. imageio.c

```

/*****
***
* Program: imageio.c
* Purpose: This source code file contains functions for dynamically allocating
* and freeing 8-bit (unsigned char) images. It also contains functions for
* reading and writing images to files in raw PGM format. This code was
written
* to be used as a teaching resource.
* Name: Michael Heath, University of South Florida
* Date: 1/7/2000
*****/

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include "imageio.h"

/*****
***
* Function: allocate_image
* Purpose: This function allocates an image. The image is an array of
pointers
* to arrays. The array of pointers will have a length of the number of rows,
* and each of these pointers will point to a separate one dimensional array
* whose length is the number of columns in the image. This scheme was used
* because it allows the image to be accessed using the syntax image[r][c]
* yet still allow the image to be any size.
* Name: Michael Heath, University of South Florida

```

```

* Date: 1/7/2000
*****
**/
unsigned char **allocate_image(int rows, int cols)
{
    unsigned char **image=NULL;
    int r, br;

    /*****
     * Allocate an array of pointers of type (unsigned char *). The array is
     * allocated to have a length of the number of rows.
     *****/

    if((image = (unsigned char **) calloc(rows, sizeof(unsigned char
*)) )==NULL){
        fprintf(stderr, "Error allocating the array of pointers in
allocate_image().\n");
        return((unsigned char **)NULL);
    }

    /*****
     * For each row, allocate an array of type (unsigned char).
     *****/

    for(r=0;r<rows;r++){
        if((image[r] = (unsigned char *) calloc(cols, sizeof(unsigned
char)))==NULL){
            fprintf(stderr, "Error allocating an array in allocate_image().\n");
            for(br=0;br<r;br++) free(image[br]);
            free(image);
            return((unsigned char **)NULL);
        }
    }

    return(image);
}

/*****
***
* Function: free_image
* Purpose: This function frees the memort that was previously allocated to
* store an image.
* Name: Michael Heath, University of South Florida
* Date: 1/7/2000
*****
**/
void free_image(unsigned char **image, int rows)
{

```

```

    int r;

/*****
 * Free each row of the image.
*****/
    for(r=0;r<rows;r++) free(image[r]);

/*****
 * Free the array of pointers.
*****/
    free(image);
}

/*****
**
 * Function: read_pgm_image
 * Purpose: This function reads in an image in raw PGM format. Because the PGM
 * format includes the number of columns and the number of rows in the image,
 * these are read from the file. Memory to store the image is allocated in
 * this
 * function. All comments in the header are discarded in the process of
 * reading
 * the image. Upon failure, this function returns 0, upon success it returns 1.
 * Name: Michael Heath, University of South Florida
 * Date: 1/7/2000
*****/
int read_pgm_image(char *infilename, unsigned char ***image, int *rows,
    int *cols)
{
    FILE *fp;
    int r;
    char buf[71];

/*****
 * Open the input image file for reading. If the file can not be opened for
 * reading return an error code of 0.
*****/
    if((fp = fopen(infilename, "r")) == NULL){
        fprintf(stderr, "Error reading the file %s in read_pgm_image().\n",
            infilename);
        return(0);
    }

```

```

/*****
 * Verify that the image is in PGM format, read in the number of columns
 * and rows in the image and scan past all of the header information.
*****/
fgets(buf, 70, fp);
if(strncmp(buf, "P5", 2) != 0){
    fprintf(stderr, "The file %s is not in PGM format in ", infilename);
    fprintf(stderr, "read_pgm_image().\n");
    fclose(fp);
    return(0);
}
do{ fgets(buf, 70, fp); }while(buf[0] == '#'); /* skip all comment lines
*/
sscanf(buf, "%d %d", cols, rows);
do{ fgets(buf, 70, fp); }while(buf[0] == '#'); /* skip all comment lines
*/

/*****
 * Allocate memory to store the image.
*****/
if((*image) = allocate_image(*rows, *cols)) == NULL) return(0);

/*****
 * Read in the image from the file, one row at a time.
*****/
for(r=0; r<(*rows); r++){
    if((*cols) != fread((*image)[r], 1, (*cols), fp)){
        fprintf(stderr, "Error reading the image data in
read_pgm_image().\n");
        fclose(fp);
        free_image((*image), *rows);
        return(0);
    }
}

fclose(fp);
return(1);
/*****
**
 * Function: write_pgm_image
 * Purpose: This function writes an image in raw PGM format. A comment can be
 * written to the header if coment != NULL. If there is a comment, it can
 * be up to 70 characters long.
 * Name: Michael Heath, University of South Florida

```

```

* Date: 1/7/2000
*****
*/
int write_pgm_image(char *outfilename, unsigned char **image, int rows,
    int cols, char *comment, int maxval)
{
    FILE *fp;
    int r;

    /*****
    * Open the output image file for writing.

    *****/
    if((fp = fopen(outfilename, "w")) == NULL){
        fprintf(stderr, "Error writing the file %s in write_pgm_image().\n",
            outfilename);
        return(0);
    }

    /*****
    * Write the header information to the PGM file.

    *****/
    fprintf(fp, "P5\n");
    if(comment != NULL)
        if(strlen(comment) <= 70) fprintf(fp, "# %s\n", comment);
    fprintf(fp, "%d %d\n", cols, rows);
    fprintf(fp, "%d\n", maxval);

    /*****
    * Write the image data to the file.

    *****/
    for(r=0;r<rows;r++){
        if(cols != fwrite(image[r], 1, cols, fp)){
            fprintf(stderr, "Error writing the image data in
write_pgm_image().\n");
            fclose(fp);
            return(0);
        }
    }

    fclose(fp);
    return(1);
}

```

### 3. imageio.h

```
/******  
***  
* Program: imageio.h  
* Purpose: This header file contains function prototypes for functions that  
* dynamically allocate and free 8-bit (unsigned char) images. It also  
contains  
* prototypes for functions that read and write images to files in raw PGM  
* format. This code was written to be used as a teaching resource.  
* Name: Michael Heath, University of South Florida  
* Date: 1/7/2000  
*****  
**/  
#ifndef _PGMIO_  
#define _PGMIO_  
  
unsigned char **allocate_image(int rows, int cols);  
  
void free_image(unsigned char **image, int rows);  
  
int read_pgm_image(char *infilename, unsigned char ***image, int *rows, int  
*cols);  
  
int write_pgm_image(char *outfilename, unsigned char **image, int rows,  
    int cols, char *comment, int maxval);  
  
int write_gray_bmp(char *outfilename, unsigned char **image, short int rows,  
short int cols);  
#endif
```

### 4. median\_serial.sh

```
#!/bin/bash  
  
CC=gcc  
EXEC=median  
SRC=median_serial.c  
COMP=imageio.c  
IN=Input_Image1.pgm  
OUT=MedianFiltered_Image_Serial.pgm  
  
if [ "$SRC" -nt "$EXEC" ]  
then  
    echo "Compiling..."
```



```
    $CC -o $EXEC $SRC $COMP
fi

./$EXEC -n 9 $IN $OUT
```

## 5. median\_parallel.sh

```
#!/bin/bash

CC=gcc
EXEC=median
SRC=median_parallel.c
COMP=imageio.c
IN1=Input_Image1.pgm
IN2=Input_Image2.pgm
IN3=Input_Image3.pgm
IN4=Input_Image4.pgm
IN5=Input_Image5.pgm
th=16

OUT=MedianFiltered_Image_Parallel.pgm

if [ "$SRC" -nt "$EXEC" ]
then
    echo "Compiling..."
    $CC -o $EXEC $SRC $COMP -fopenmp
fi

./$EXEC -n 9 $IN1 $IN2 $IN3 $IN4 $IN5 $OUT $th
```

## **B. Hpc-class's compute node's architecture:**

Architecture: x86\_64

CPU op-mode(s): 32-bit, 64-bit

Byte Order: Little Endian

CPU(s): 16

On-line CPU(s) list: 0-15

Thread(s) per core: 1

Core(s) per socket: 8

Socket(s): 2

NUMA node(s): 2

Vendor ID: GenuineIntel

CPU family: 6

Model: 45

Model name: Intel(R) Xeon(R) CPU E5-2650 0 @ 2.00GHz

Stepping: 7

CPU MHz: 1217.895

CPU max MHz: 2800.0000

CPU min MHz: 1200.0000

BogoMIPS: 4000.38

Virtualization: VT-x

L1d cache: 32K

L1i cache: 32K

L2 cache: 256K

L3 cache: 20480K

NUMA node0 CPU(s): 0-7

NUMA node1 CPU(s): 8-15

available: 2 nodes (0-1)

node 0 cpus: 0 1 2 3 4 5 6 7

node 0 size: 32735 MB

node 0 free: 31611 MB

node 1 cpus: 8 9 10 11 12 13 14 15

node 1 size: 32768 MB

node 1 free: 31639 MB

node distances:

```
node 0 1
0: 10 21
1: 21 10
```

### **C. Memory Usage:**

As I could not find a very large data set of .pgm images, I had to download large resolution .jpg images and convert them online to .pgm image format. Due to this, I was unable to get large chunk of data set which would have increased my data-set size.

As for now, I have used less than 1 gigabyte of memory including the five input images, the five output images and some additional parameters.

## REFERENCES

---

- [1] <https://www.mathworks.com/help/images/ref/medfilt2.html>
- [2] [https://www.cs.auckland.ac.nz/courses/compsci373s1c/PatricesLectures/Image%20Filtering\\_2up.pdf](https://www.cs.auckland.ac.nz/courses/compsci373s1c/PatricesLectures/Image%20Filtering_2up.pdf)
- [3] An Improved Median Filtering Algorithm for Image Noise Reduction - Youlian Zhu, Cheng Huang  
  
([https://ac.els-cdn.com/S1875389212005494/1-s2.0-S1875389212005494-main.pdf?\\_tid=5d583886-669c-4c49-a8df-cec956c8c851&acdnat=1540960253\\_d76b3e6405bd110c4c0f65251d38deaf](https://ac.els-cdn.com/S1875389212005494/1-s2.0-S1875389212005494-main.pdf?_tid=5d583886-669c-4c49-a8df-cec956c8c851&acdnat=1540960253_d76b3e6405bd110c4c0f65251d38deaf))
- [4] <https://www.ece.rice.edu/~wakin/images/>
- [5] <https://www.cs.cmu.edu/~eugene/teach/algs00a/progs/>
- [6] <https://www.youtube.com/watch?v=APkdYObUJRU>
- [7] Data-set source: <https://unsplash.com/search/photos/4k>  
Converted .jpg to .pgm image from: <https://convertio.co/jpg-pgm/>
- [8] <https://www.openmp.org/wp-content/uploads/openmp-4.5.pdf>
- [9] <https://www.openmp.org/wp-content/uploads/openmp-examples-4.5.0.pdf>